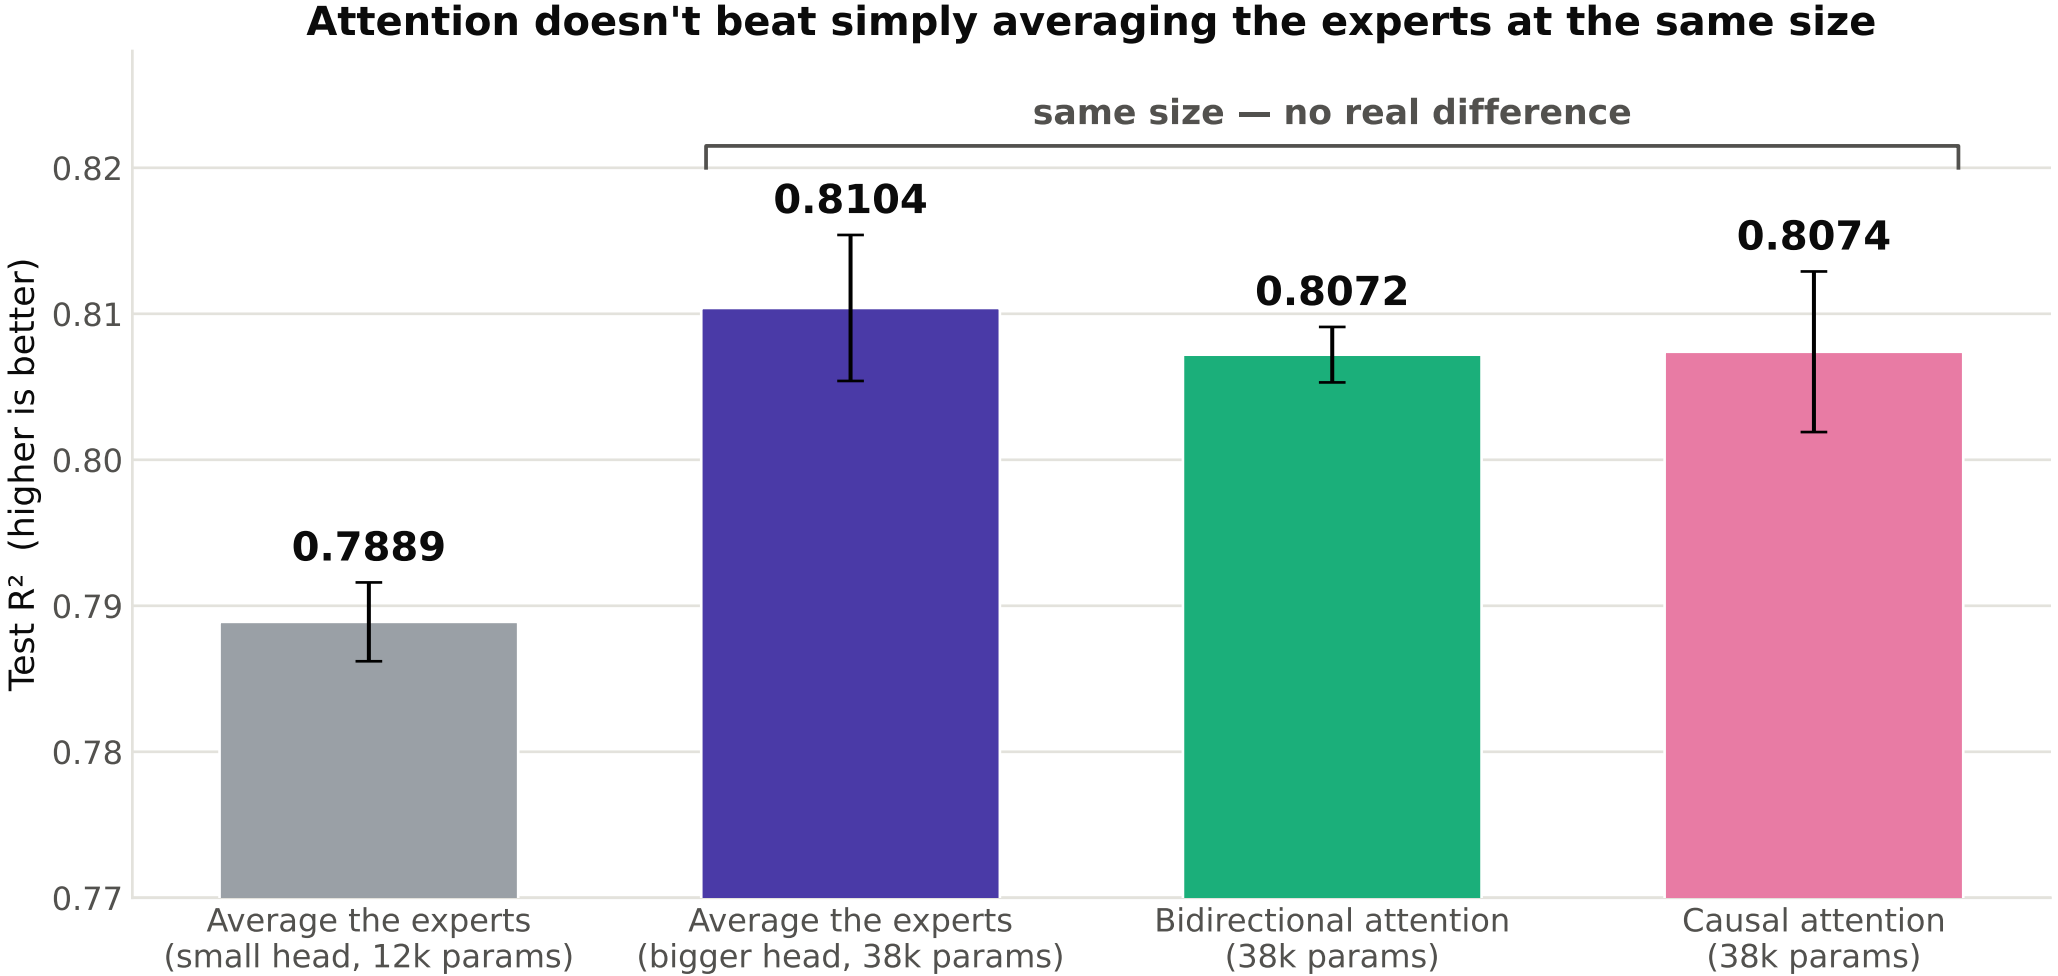
Attention doesn't beat simply averaging the experts at the same size
same size — no real difference
Test R² (higher is better)
0.82
0.81
0.80
0.79
0.78
0.77
0.7889
0.8104
0.8072
0.8074
Average the experts (small head, 12k params)
Average the experts (bigger head, 38k params)
Bidirectional attention (38k params)
Causal attention (38k params)